

# CSCE 3600

## Principles of Systems Programming

### Interprocess Communication Part 1

University of North Texas



**CSE**

## Interprocess Communication



## Interprocess Communication

- Processes may be independent or cooperating
  - An **independent process** cannot affect or be affected by the execution of another process
    - A process that does not share data with another process is independent
  - A **cooperating process** can affect or be affected by the execution of another process
    - A process that shares data with another process is a cooperating process
    - Cooperating processes require **interprocess communication (IPC)**

3



## Cooperating Processes



- Why provide an environment for cooperating processes?
  - **Information sharing**
    - Cooperating processes can share information (such as access to the same files) among multiple processes, but a mechanism is required for parallel access
  - **Modularity**
    - Involves dividing complicated tasks into smaller subtasks, which can be completed by different *cooperating* processes
  - **Computation speedup**
    - Subtasks of a single task can be performed in parallel using cooperating processes
  - **Convenience**
    - Different tasks completed/managed by cooperating processes

4



## Methods of Cooperation

- Cooperation by sharing
  - Use shared data such as memory, variables, files, databases, etc. mutually exclusive
  - Critical section used to provide data integrity with mutually exclusive writing to prevent inconsistent data



- Cooperation by communication
  - Cooperation using messages (**message passing**)
  - May lead to deadlock if each process waiting for a message from the other to perform an operation
  - Starvation also possible if process never receives a message

We will focus on message passing in this class

5



## Message Passing

- A mechanism is needed for processes to *communicate* and *synchronize* their actions
- In a **message system**, processes communicate with each other without resorting to shared variables
- IPC facility provides two primitive operations for fixed or variable-sized message passing: **send** and **receive**
- If two processes wish to communicate, they need to
  - Establish a communication link between them
    - Logical (e.g., logical properties) or physical (e.g., shared memory, hardware bus)
  - Exchange messages via send/receive

6



## Synchronizing Messages

- Message passing may be either **blocking** or **non-blocking**
  - Blocking is considered to be **synchronous**
    - **Send** sender blocked until message received
    - **Receive** receiver blocks until a message is available
  - Non-Blocking is considered to be **asynchronous**
    - **Send** sender resumes operation immediately after sending (i.e., send and continue)
    - **Receive** receiver returns immediately with either a valid or null message

7



## Buffering

- A **message queue** can be implemented in one of three ways
  - **Zero capacity**
    - No messages may be queued within the link
      - Requires sender to wait until receiver retrieves message
  - **Bounded capacity**
    - Link has finite number of message buffers
      - If no buffers are available, then sender must wait if the link is full
  - **Unbounded capacity**
    - Link has unlimited buffer space, so sender never needs to wait

8



## IPC Methods

---

- We will explore the following IPC methods
  - Signaling
    - As a limited form of IPC, a signal is essentially an asynchronous notification sent to a process in order to notify it of an event that occurred
  - Files
    - A file is a **durable** block of arbitrary information, or resource for storing information
  - Pipes
    - A pipe is a synchronized, interprocess byte stream
  - Sockets
    - Sockets provide point-to-point, two-way communication between two processes, even processes on different systems

9

**CSE**

---

## Signal Handling

10



## What Are Signals?

- A **signal** is a software interrupt
  - Notification of an event
    - A way to communicate information to a process about the state of other processes, the OS, and hardware so that the process can take appropriate action
  - Can change the flow of the program
    - When a signal is delivered to a process, process will stop what it's doing – and either **handle** or **ignore** signal
  - Signals can be delivered in an unpredictable manner
    - Originate outside of currently executing process
    - **Asynchronous** events due to external trigger at the hardware or OS level – causes a context switch!

11



## What Are Signals?

- Every signal has a name that starts with SIG, a value, a default action, and a description
  - See [man 7 signal](#)

Signal	Value	Action	Comment
<b>SIGHUP</b>	1	Term	Hangup detected on controlling terminal or death of controlling process
<b>SIGINT</b>	2	Term	Interrupt from keyboard
<b>SIGQUIT</b>	3	Core	Quit from keyboard
<b>SIGILL</b>	4	Core	Illegal Instruction
<b>SIGABRT</b>	6	Core	Abort signal from <b>abort(3)</b>
<b>SIGFPE</b>	8	Core	Floating point exception
<b>SIGKILL</b>	9	Term	Kill signal
<b>SIGSEGV</b>	11	Core	Invalid memory reference
<b>SIGPIPE</b>	13	Term	Broken pipe: write to pipe with no readers
<b>SIGALRM</b>	14	Term	Timer signal from <b>alarm(2)</b>
<b>SIGTERM</b>	15	Term	Termination signal
<b>SIGUSR1</b>	30,10,16	Term	User-defined signal 1
<b>SIGUSR2</b>	31,12,17	Term	User-defined signal 2
<b>SIGCHLD</b>	20,17,18	Ign	Child stopped or terminated
<b>SIGCONT</b>	19,18,25	Cont	Continue if stopped
<b>SIGSTOP</b>	17,19,23	Stop	Stop process
<b>SIGTSTP</b>	18,20,24	Stop	Stop typed at tty
<b>SIGTTIN</b>	21,21,26	Stop	tty input for background process
<b>SIGTTOU</b>	22,22,27	Stop	tty output for background process

Defined in  
[sys/signal.h](#)

12



## Default Action of Signals

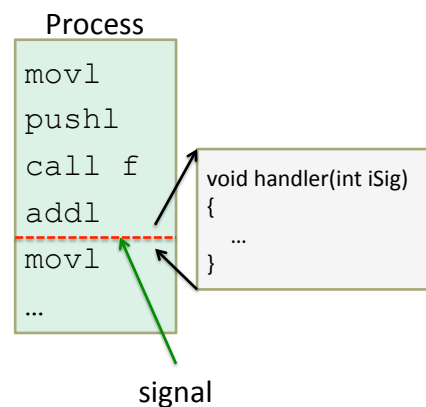
- Each signal has a default action
  - Term      the process will **terminate**
  - Core      the process will terminate and produce a **core dump** file that traces the process state at the time of termination
  - Ign      the process will **ignore** the signal
  - Stop      the process will **stop**, like Ctrl-Z
  - Cont      the process will **continue** from being stopped

13



## Flow of Signals

1. Event gains attention of OS
2. OS stops process execution immediately
  - Sends it a signal
3. **Signal Handler** then executes to completion
4. Process execution resumes where it left off



14



## Signal Events

- In the context of **terminal signaling**, programs can stop, start, and terminate
  - `Ctrl-C` is the same as sending `SIGINT` (2) signal
    - Default handler exits process
  - `Ctrl-Z` is the same as sending a `SIGSTOP` (20) signal
    - Default handler suspends process
  - `Ctrl-\` is the same as sending a `SIGQUIT` (3) signal
    - Default handler exits process
  - Typing `fg` or `bg` at the terminal is the same as sending a `SIGCONT` (18) signal to bring or send a process to the **foreground** or **background**, respectively

15



## Signal Events

- Signals are notification of events
  - We can **inject** signals, such as `Ctrl-C`, to gain the attention of the OS and stop the process by sending a `SIGINT` signal
  - But some are done **internally**, such as when a process makes an illegal memory reference
    - Event gains attention of the OS
    - OS stops process execution immediately, sending it a `SIGSEGV` (11) signal
    - Signal handler for `SIGSEGV` signal executes to completion
      - Default signal handler for `SIGSEGV` signal prints “segmentation fault” and exits process

16





## Signal Handling

- Each signal type has a default handler
  - Most default handlers exit the process
- A program can install its own handler for signals of **almost** any type
  - Cannot install its own handler for the following signals
    - SIGKILL (9)
      - Default handler exits the process
      - Catchable termination signal is SIGTERM (15)
    - SIGSTOP (19)
      - Default handler suspends the process
      - Can resume the process with signal SIGCONT (18)
      - Catchable suspension signal is SIGTSTP (20)

17



## Installing a Signal Handler

- `sighandler_t signal(int iSig, sighandler_t pfHandler);`
  - Installs `pfHandler` as the handler of signals for type `iSig`
  - `pfHandler` is a function pointer
 

```
typedef void (* sighandler_t) (int);
```
  - Returns the old handler on success; `SIGERR` on error
  - After call, `pfHandler` is invoked whenever process receives a signal of type `iSig`

18



## Signal Handler Example

```
#include <stdio.h>
#include <assert.h>
#include <signal.h>

void myHandler(int iSig) {
    printf("In myHandler with argument %d\n", iSig);
}

int main() {
    void (*pfRet)(int) = signal(SIGINT, myHandler);
    assert(pfRet != SIG_ERR);

    printf("Entering an infinite loop\n");
    while (1) {
        printf(".");
    }
    return 0;
}
```

19



## Predefined Signal Handlers

- Can install predefined signal handler `SIG_IGN` to ignore signals  
`void (*pfRet)(int) = signal(SIGINT, SIG_IGN);`
  - Subsequently, process will ignore `SIGINT` (2) signals
- Can install predefined signal handler `SIG_DFL` to restore default signal handler  
`void (*pfRet)(int) = signal(SIGINT, myHandler);`  
...  
`pfRet = signal(SIGINT, SIG_DFL);`
  - Subsequently, process will handle `SIGINT` (2) signals using the default handler for `SIGINT` (2) signals

20



## Sending Signals via Command

- `kill -signal pid`
  - Send signal of type `signal` to process with ID `pid`
  - Specify signal type name (`-SIGINT`) or number (`-2`)
  - Examples:
 

<code>kill -2 1234</code>	}	Same as typing <code>Ctrl-C</code> if process
<code>kill -SIGINT 1234</code>		1234 is running in foreground
- `killall loop`
  - Send a signal to `all` processes named `loop` to “terminate”
  - This will actually send the signal `SIGTERM`, whose purpose is to communicate a termination request to the given process, which does not necessarily have to terminate

21



## Sending Signals via System Call

- `int raise(int iSig);`
  - Instructs OS to send signal of type `iSig` to current process
  - Returns 0 to indicate success; non-0 to indicate failure
  - Example:
 

```
int retVal = raise(SIGINT); // process commits suicide
assert(retVal != 0);       // should not get here
```
- `int kill(pid_t iPid, int iSig);`
  - Sends `iSig` signal to process with ID `iPid`
  - Equivalent to `raise(iSig)` when `iPid` is ID of current process
  - Example:
 

```
pid_t iPid = getpid(); // process gets its pid
kill(iPid, SIGINT);    // sends itself a SIGINT signal, commits suicide
```

22



## pause ( ) Function

- `int pause( );`
  - Suspends the calling process, without wasting resources, until some kind of signal is caught
  - Signal action can be the execution of a signal handler function or process termination
  - Only returns `-1` when a signal was caught and the signal-catching function returned

23



## pause ( ) Example

```
#include <stdio.h>
#include <unistd.h>
#include <signal.h>

void sig_usr( int signo );    /* handles two signals */

int main() {
    int i = 0;
    if( signal( SIGUSR1, sig_usr ) == SIG_ERR )
        printf( "Cannot catch SIGUSR1\n" );
    if( signal( SIGUSR2, sig_usr ) == SIG_ERR )
        printf( "Cannot catch SIGUSR2\n" );
    :
    while(1) {
        printf( "%2d\n", i );
        pause();
        /* pause until signal handler
         * has processed signal */
        i++;
    }

    return 0;
}
```

24



## pause () Example

```
/* argument is signal number */
void sig_usr( int signo )
{
    if( signo == SIGUSR1 )
        printf( "Received SIGUSR1\n" );
    else if( signo == SIGUSR2 )
        printf( "Received SIGUSR2\n" );
    else
        printf( "Error: received signal %d\n", signo );

    return;
}
```

```
$ sig_examp &
[1] 4720
0
$ kill -USR1 4720
Received SIGUSR1
1
$ kill -USR2 4720
Received SIGUSR2
2
$ kill 4720 /* send SIGTERM */
[1] + Terminated sig_examp &
$
```

25



# CSE

## Files

26



## Overview of Files

- Data communication with a C program and the outside world is performed through files
- Files are a non-volatile way to store data by means of such media as tape, CD-ROM, floppy disk, ZIP disk, hard drive, etc.
- C (just like the Linux operating system) considers all process communication media to be **files**
  - An ordinary file on a disk is considered to be a file
  - So is the keyboard, the screen, parallel ports, and serial ports

27



## Linux Files

- A Linux file is a sequence of  $m$  bytes
  - $B_0, B_1, \dots, B_k, \dots, B_{m-1}$
- All I/O devices are represented as files
 

<code>/dev/sda2</code>	( <code>/usr</code> disk partition)
<code>/dev/tty2</code>	(terminal)
- Even the kernel is represented as a file
 

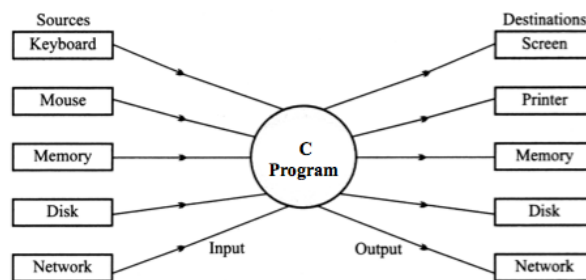
<code>/dev/kmem</code>	(kernel memory image)
<code>/proc</code>	(kernel data structures)

28



## I/O and Data Movement

- Input and output share common property of **unidirectional** movement of data and support to **sequential access** to data
  - The flow of data into a program (input) may come from different devices such as keyboard, mouse, memory, disk, network, or another program
  - The flow of data out of a program (output) may go to the screen, printer, memory, disk, network, another program



29



## Some Linux File Types (Review)

- Regular file
  - Binary or text file (Linux does not know the difference)
- Directory file
  - A file that contains the names and locations of other files
- Character special and block special files
  - Terminals (character special) and disks (block special)
- FIFO (named pipe)
  - A file type used for interprocess communication to exchange data between unrelated processes
- Socket
  - A file type used for network communication between processes

30



## Communication through Files

- Storing and manipulating data using files is known as **file processing**
- Programs access files through basic file operations
  - Open a file
  - Read data from a file
  - Write data to a file
  - Close a file
 } these “process” a file
- Text files store all data types as character bytes
  - The way a program reads the data (either text or binary mode) determines how the data is interpreted
  - Example: 12 z rti 456.79 room

31



## Standard I/O Functions

- Examples of standard I/O functions:
  - Opening and closing files (`fopen` and `fclose`)
  - Reading and writing bytes (`fread` and `fwrite`)
  - Reading and writing text lines (`fgets` and `fputs`)
  - Formatted reading and writing (`fscanf` and `fprintf`)

```
FILE* fp;
fp = fopen("In.file", "rw");
fscanf(fp, .....);
fprintf(fp, .....);
fread(....., fp);
fwrite(....., fp);
```

32





## Standard I/O Streams

- Standard I/O models open files as *streams*
  - Abstraction for a file descriptor and a buffer in memory.
- Three files are automatically opened each time a C program is run and automatically closed when a C program ends:
  - `stdin` (standard input, default source = keyboard)
  - `stdout` (standard output, default destination = screen)
  - `stderr` (standard error, default destination = screen)

```
#include <stdio.h>
extern FILE *stdin; /* standard input (descriptor 0) */
extern FILE *stdout; /* standard output (descriptor 1) */
extern FILE *stderr; /* standard error (descriptor 2) */

int main() {
    fprintf(stdout, "Hello, world\n");
}
```

streams


file descriptors

33



## Buffering in Standard I/O

- Standard I/O functions use buffered I/O



34



## Standard I/O Buffering Example

- You can see this buffering using the Linux `strace` program:

```
#include <stdio.h>

int main()
{
    printf("h");
    printf("e");
    printf("l");
    printf("l");
    printf("o");
    printf("\n");
    fflush(stdout);
    return 0;
}
```

```
linux> strace ./hello
execve("./hello", ["hello"], [/ * ... */]).
...
write(1, "hello\n", 6...)           = 6
...
_exit(0)                           = ?
```

35



## File Descriptors

- Each open file is associated with an **open file description**
  - An OS internal record of how a process or group of processes are *currently* accessing a file that includes
    - File offset** indicates byte position where next I/O operation begins
    - File status** indicates append mode/not, blocking/non-blocking, etc.
    - File access** mode indicates whether file can be read or written
- Each process has a logical array of references to open file descriptions
  - Logical indices into this array are **file descriptors** used to identify the files for I/O operations
    - A file descriptor does not describe a file – it's just a number ephemerally associated with a particular open file description
  - List open files: `ls -l /proc/<pid>/fd`

36



## File Descriptors


- Other descriptors are assigned by system calls to open/create files, create pipes, or bind to devices or network sockets
  - E.g., `pipe()`, `socket()`, `open()`, `creat()`
- A common set of system calls operate on open file descriptors independent of their underlying types
  - E.g., `read()`, `write()`, `dup()`, `dup2()`, `close()`

37



## Files as Pipes for Data

- Files act as *pipes* to bring a stream of data into the program and send a stream of data out of the program



- The program reads data from `stdin` and writes data to `stdout` as if they were *ordinary* files

38



## One Byte at a Time

- Data is read or written one byte at a time from a file until the end of the file is reached or until the file is closed
- The file system uses a pointer to keep track of the next byte to read or to write

Smith	Jack	1045.76	Manager	15
Hanson	Susan	98.62	Operator	7
Jones	Nancy	790.25	Administrator	10
Doe	Carl	526.71	Technician	12

- In this file, the program has read the data as far as the letter 'J'. The next character to be read is 'a'

39



## Functions Using stdin and stdout

- Some standard C functions **implicitly** use `stdin` and `stdout`

```
scanf("%d", &aNumber);
aSymbol = getchar();
gets(theBuffer); // High security risk
printf("Average: %5.2f", theAverage);
putchar(aCharacter);
puts(aPhrase);
```

40



## Functions Using stdin and stdout

---

- Other functions need to have `stdin` and `stdout` specified as the file descriptor

```
fscanf(stdin, "%d", &aNumber);  
aSymbol = fgetc(stdin);  
fgets(theBuffer, sizeof(theBuffer), stdin);  
fprintf(stdout, "Average: %f", theAverage);  
fputc(aCharacter, stdout);  
puts(aPhrase, stdout);
```

41

**CSE**

## Pipes

42



## What is a Pipe?

- A **pipe** is a simple, **synchronized** way of passing information between processes
  - A pipe is a special file/buffer that stores a limited amount of data in a FIFO (i.e., sequential) manner
  - Pipes are commonly used from within shells to connect `stdout` of one utility to `stdin` of another
- Pipes provide a basic level of synchronization
  - If a process tries to write to a **full pipe**, it is **blocked** until the reader process consumes some data
  - If a process tries to read from an **empty pipe**, it is **blocked** until the writer produces some data
- Data is written to and read from the pipe using the unbuffered system calls `write()` and `read()`

43



## Two Types of Pipes

- **Unnamed pipes**
  - Not associated with any file
  - Can only be used with related processes, such as a parent and child process
  - Exist only as long as using processes are not terminated and support **unidirectional** communication
  - Created using the `pipe()` system call
- **Named pipes, or FIFOs**
  - Associated with a file and can be used with unrelated processes
  - Has a directory entry with file access permissions
  - Can support **bidirectional** communication
  - Created using the `mknod()` or `mkfifo()` system calls

44



## Unnamed Pipes

- The `pipe()` system call creates an unnamed pipe and opens it for reading and writing
- General syntax

```
int pipe(int fd[2]);
```

- If successful, it will return two integer file descriptors in `fd[0]` and `fd[1]`
- `fd` must be an integer array of size 2
- The file descriptor in `fd[0]` is associated with the read end of the pipe and `fd[1]` is associated with the write end of the pipe
- Returns -1 if an error occurred

45




## Pipes



File Descriptor Table

0	stdin
1	stdout
2	stderr
3	fd[0]
4	fd[1]
5	...



46



## Writing to a Pipe

```
#include <unistd.h>
ssize_t write(int fd, const void *buf, size_t count);
```

- Appends up to `count` bytes from the end of the pipe referenced by the file descriptor from the buffer
- **Atomicity** guaranteed for requests with size typically around 4096 bytes or less
  - See `PIPE_BUF` in `/usr/include/linux/limits.h` for block size
- If `write()` done for pipe not open for reading by any process
  - `SIGPIPE` signal generated to signify a broken pipe with `errno` set to `EPIPE` (broken pipe)
- See [man 2 write](#) for more information on this system call

47



## Reading from a Pipe

```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t count);
```

- Attempts to read up to `count` bytes from end of the pipe referenced by the file descriptor from the buffer in a FIFO manner
  - Returns number of bytes read from the buffer
- If `read()` done for pipe not open for writing by any process, 0 is returned
- If `read()` done for empty pipe open for writing by another process, the process sleeps until the input becomes available
- See [man 2 read](#) for more information on this system call

48






## Creating and Using Pipes

- Created using `pipe()`:  

```
int filedes[2];
pipe(filedes);
.
.
.
write(filedes[1], buf, count);
read(filedes[0], buf, count);
```
- Pipes are anonymous
  - There is no name in the file system
- But then how do two processes share a pipe?



49



## Unnamed Pipes

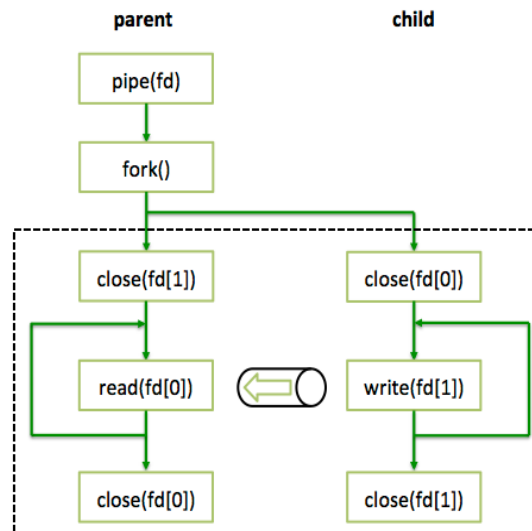
- Since access to an unnamed pipe is through the file descriptor mechanism, **only the process that created the pipe and its descendants may use the pipe**
- The typical sequence of opening unnamed pipes
  - The parent process creates an unnamed pipe
    - It is crucial that this be done **before** forking
  - The parent process forks
  - The writer process closes the read end of the pipe
  - The reader process closes the write end of the pipe
  - The processes communicate by using `write()` and `read()`
  - Each process closes its active pipe-end when finished

May be reversed

50



## Pipe Communication Scheme




51



## Sharing a Pipe

```

int filedes[2];
pipe(filedes);
pid = fork();
if (pid == 0) {
    close(filedes[1]);
    // child now reads
} else {
    close(filedes[0]);
    // parent now writes
}
  
```



- The `fork()` system call duplicates parent's file descriptors
  - Parent and child must close unused descriptors – necessary for correct use of pipes

52



## Pipe Example (1)

```
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#define READ 0 /* The index of the read end of the pipe */
#define WRITE 1 /* The index of the write end of the pipe */

char* phrase = "This goes in the pipe";

int main () {
    int fd[2], bytesRead;
    char message[100]; /* Parent process' message buffer */

    pipe(fd); /* Create unnamed pipe */
    if (fork() == 0) /* Child, writer */
    {
        close(fd[READ]); /* Close unused end */
        write(fd[WRITE], phrase, strlen(phrase) + 1); /* Include NULL */
        close(fd[WRITE]); /* Close used end */
    }
    else /* Parent, reader */
    {
        close(fd[WRITE]); /* Close unused end */
        bytesRead = read(fd[READ], message, 100);
        printf("Parent just read %i bytes: %s\n", bytesRead, message);
        close(fd[READ]); /* Close used end */
    }
}
```

53



## Pipe Example (2)

```
/* This program will demonstrate what happens if a read takes
   place with a pipe whose write end is closed, and vice versa */
#include <stdio.h>
#include <signal.h>
#include <unistd.h>
#include <stdlib.h>
#include <string.h>
#define READ 0 /* The index of the read end of the pipe */
#define WRITE 1 /* The index of the write end of the pipe */
char* phrase = "Another pipe example end closed";
void signal_catcher(int);
int main() {
    int fd[2], bytesWritten = 0, bytesRead;
    char message[100]; /* Parent process' message buffer */
    signal(SIGPIPE, signal_catcher);
    pipe(fd); /* Create pipe */
    close(fd[WRITE]); /* Close used end */
    printf("About to read from pipe\n");
    bytesRead = read(fd[READ], message, 100);
    printf("%i bytes were read with write closed\n", bytesRead);
    close(fd[READ]); /* Close used end */
    pipe(fd); /* Recreate unnamed pipe */
    close(fd[READ]); /* Close unused end */
    printf("About to write to pipe\n");
    bytesWritten = write(fd[WRITE], phrase, strlen(phrase) + 1);
    printf("%i bytes were written with read end closed\n", bytesWritten);
    close(fd[WRITE]);
}
void signal_catcher(int theSig) {
    printf("A SIGPIPE (%i) has been caught\n", theSig);
}
```

54



## dup ( ) System Call

```
#include <unistd.h>
int dup(int oldfd);
```

- Creates a copy of the `oldfd` file descriptor with the same open pipe, file pointer, and access mode in common with the original file descriptor that is set to remain open across the `exec` family of system calls
- Returns the lowest `unused` file descriptor, -1 if error
  - Problems if file descriptor returned is not one you are expecting!
- A useful system call to convert a stream to a file descriptor

```
int fileno(FILE *fp);
```

Note that `dup()` refers to a file descriptor, not a stream!

55



## dup ( ) Example

```
int fd[2];
pipe(fd);
close(fileno(stdout));
dup(fd[1]);
```

Note that no error checking is done in this example

- Since 1 is the lowest `fd` available, the write end of the pipe is duplicated at fd 1 (`stdout`)
  - Now any data written to `stdout` will be written to the pipe
- But you are taking a chance that the file descriptor that will be returned by `dup()` is what you want
  - The process may be interrupted between the `close()` and the `dup()`

56



## dup2 ( ) System Call

```
#include <unistd.h>
int dup2(int oldfd, int newfd);
```

- A better alternative to `dup()` in that it makes the `newfd` file descriptor as the copy of the old file descriptor **atomically**, closing `newfd` first if necessary
  - There is no time lapse between closing `newfd` and duplicating `oldfd` into its spot
- Both `oldfd` and `newfd` now refer to the same file

57



## dup ( ) and dup2 ( ) Example

```
/* This program demonstrates the dup and dup2 system calls.
   You must have a file present in the directory called "test.txt".
   It may be empty or have stuff in it doesn't matter. */
#include <stdio.h>
#include <sys/types.h>
#include <fcntl.h>
#include <sys/file.h>
int main() {
    int fd1, fd2, fd3;
    fd1 = open("test.txt", O_RDWR | O_TRUNC);
    printf("fd1 = %i\n", fd1);
    write(fd1, "what's", 6);
    fd2 = dup(fd1); /* make a copy of fd1 */
    printf("fd2 = %i\n", fd2);
    write(fd2, " up", 3);
    close(0); /* close standard input */
    fd3 = dup(fd1); /* make another copy of fd1 */
    printf("fd3 = %i\n", fd3);
    write(0, " doc", 4); /* because 0 was the smallest file descriptor */
                        /* and now belongs to fd3 */
    dup2(3,2); /* duplicate channel 3 to channel 2 */
    write(2, "?\n", 2);
}
```

58





## Implement Command-Line Pipe

```

/* Modeling the command-line command: ls -al | wc using pipes */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
enum { READ, WRITE };
int main() {
    int fd[2];
    if (pipe(fd) == -1) { /* generate the pipe */
        perror("Pipe");
        exit(1);
    }
    switch (fork()) {
        case -1: perror("Fork");
                exit(2);
        case 0: /* in child */
                dup2(fd[WRITE], fileno(stdout));
                close(fd[READ]);
                close(fd[WRITE]);
                execl("/bin/ls", "ls", "-al", (char *)0);
                exit(3);
        default: /* in parent */
                dup2(fd[READ], fileno(stdin));
                close(fd[READ]);
                close(fd[WRITE]);
                execl("/usr/bin/wc", "wc", (char *)0);
                exit(4);
    }
    return 0;
}

```

61



## Implement Redirection

- When a process forks, the child inherits a copy of its parent's file descriptors
- When a process execs, all non-close-on-exec file descriptors remain unaffected
  - This includes `stdin`, `stdout`, and `stderr`
- To implement redirection, the shell simply does the following
  - The parent shell forks then waits for the child shell to terminate

62



## Implement Redirection

```

/* The program demonstrates implementing redirection.
   To run: ./a.out <output filename> <command>
*/
#include <stdio.h>
#include <sys/file.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
int main(int argc, char* argv[])
{
    int fd;
    /* Open file for redirection */
    fd = open(argv[1], O_CREAT | O_TRUNC | O_WRONLY, 0600);
    dup2(fd, 1); /* Duplicate descriptor to standard output */
    close(fd); /* Close original descriptor to save descriptor space */
    execvp(argv[2], &argv[2]); /* Invoke program; will inherit stdout */
    perror("main"); /* Should never execute */
}

```

63



## Implement Redirection

- The child shell opens the file, say `ls.out`, creating it or truncating it as necessary
- The child shell then
  - Duplicates the file descriptor of `ls.out` to the standard output file descriptor (`fd 1`)
  - Closes the original file descriptor of `ls.out`
    - All standard output is therefore directed to `ls.out`
  - The child shell then execs the `ls` utility
  - Since the file descriptors are inherited during an exec, all `stdout` of `ls` goes to `ls.out`
- When the child process terminates, the parent resumes
- The parent's file descriptors are unaffected by the child's action as each process maintains its own descriptor table

64





## Implement Pipe and Redirection

```

/* Equivalent to "sort < file2 | uniq */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
int main() {
    int fd[2];
    FILE *fp = fopen("file2", "r");
    dup2(fileno(fp), fileno(stdin));
    fclose(fp);
    pipe(fd);
    if (fork() == 0) {
        dup2(fd[1], fileno(stdout));
        close(fd[0]);
        close(fd[1]);
        execl("/usr/bin/sort", "sort", (char*) 0);
        exit(2);
    } else {
        dup2(fd[0], fileno(stdin));
        close(fd[0]);
        close(fd[1]);
        execl("/usr/bin/uniq", "uniq", (char*) 0);
        exit(3);
    }
    return 0;
}

```

65