

CSCE 3600

Principles of Systems Programming

Processes

University of North Texas



CSE

Process Management



What is a Process?

- A **program** is a *passive* set of instructions stored on a secondary storage device, such as a disk
- A **process** is an *active* execution of a program stored in memory
 - Program becomes process when loaded into memory
- Processes can create sub-processes to execute concurrently
- The execution of a process must progress in a sequential fashion
 - The CPU executes one instruction of the process after another until the process completes, or other event occurs

3



More Precise Definition

- A process is the **context** (i.e., information and data) maintained for an executing program
 - A process needs certain resources, including CPU time, memory, files, and I/O devices, to accomplish its task
- Intuitively, a process is the abstraction of a physical processor
 - Exists because it is difficult for the OS to otherwise coordinate many concurrent activities, such as incoming network data, multiple users, etc.

4



Process Management

- The OS is responsible for the following activities
 - Process creation and deletion
 - Process suspension and resumption
 - Provision of mechanisms for:
 - Process synchronization
 - Process communication
 - Deadlock handling
- How does OS correctly run multiple processes concurrently?
 - What kind of information must be kept?
 - What does OS have to do to run processes correctly?
 - OS must be able to distinguish among different processes
 - Multiple programs may be loaded into memory at same time
 - Each process is assigned a unique, non-negative integral **process ID**, or PID

5



Process Identifiers

- Special processes with well-known process IDs
 - **swapper / sched**
 - PID 0, as part of the kernel, a system process responsible for memory management
 - **init** replaced by **systemd** in many Linux distributions
 - PID 1, a continually-running daemon process (i.e., one that runs in the background) responsible for starting up and shutting down the system
 - Invoked by the kernel at the end of the bootstrap procedure

Check out the **ps tree** command that prints a tree of the processes

6




Process Context

7



Process Context

- When a program loaded into memory, it is organized into the following segments of memory
 - **text** contains the actual program code, or executable instructions
 - **data** contains global and static variables initialized at runtime
 - **heap** contains dynamic memory allocated at runtime
 - **stack** contains return addresses, function parameters, and variables



This is known as **User Level Context**

8

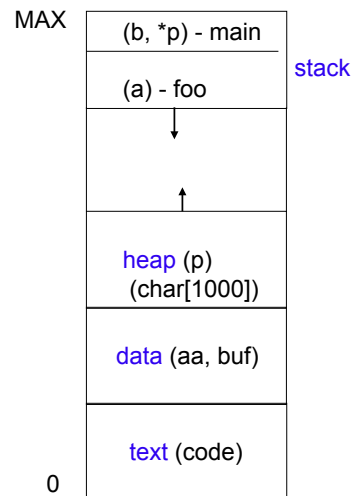


User Level Context

```

...
int aa;
char buf[1000];
void foo() {
    int a;
    ...
}
main() {
    int b;
    char *p;
    p = new char[1000];
    foo();
}

```




9



Process Context

- Is the *user level context* sufficient?
 - Does it contain all the states necessary to run a program?
 - Only if the system runs through one program at a time
 - The OS typically needs to switch back and forth between programs – processes must be “swapped” in and out from getting to use the CPU
- R_1 in P_1 is incorrect... why? How to make it right?
 - Save R_0 in P_1 before switching
 - Restore R_0 in P_1 when switching from P_2 to P_1
- Registers should be a part of process context
 - This is called **Register Context**



10



Register Context

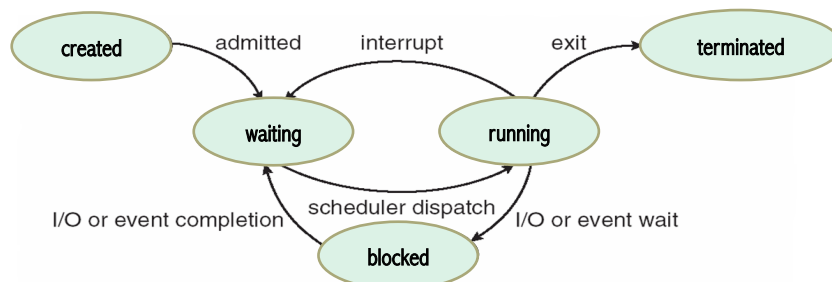
- We need to save everything that we need to independently run a process
 - Program Counter (PC)
 - Address of next instruction to be executed (may be in kernel or user memory space of this process)
 - Processor Status Register
 - Contains the hardware status at the time of preemption – contents and format are hardware dependent
 - Stack Pointer (SP)
 - Points to the top of the kernel or user stack, depending on the mode of operation at the time of preemption
 - General-Purpose Registers
 - Hardware dependent, R_0, R_1, R_2, \dots

11



Process State

- A process executes according to the following state machine:
 - created also “new”, initial state when created
 - waiting also “ready”, awaiting to be scheduled for execution
 - running actively executing instructions on the CPU
 - blocked unable to continue without event occurring, e.g., I/O
 - terminated no longer running due to completion or being killed



12



Process Context

- User level context
 - Code, data, stack, heap
- Register context
 - R0, R1, ..., PC, SP, etc.
- What else is needed?
 - OS resources
 - Open files, signal related data structures, etc.

While a program is executing, the process can be uniquely identified by a number of elements, including:

- Identifier (PID)
- State
- Priority
- Program counter
- Memory pointers
- Context data
- I/O status information
- Accounting information

To run a process correctly, the process instructions must be executed within the process context!

13



Process Context

- Where is the *process context* stored?
 - User level context is in memory
 - Other context information is stored in a data structure called **process control block**
 - Contains other information that the OS needs to manage the process
 - Process status (running, waiting, etc.)
 - Process priority
 - ...
 - The OS has a **process control block table**
 - For each process, there is one entry in the table


14



Process Control Block (PCB)

Information associated with each process (also called **task control block**)

- Process state
 - Running, waiting, etc.
- Program counter
 - Location of instruction to execute next
- CPU registers
 - Contents of all process-centric registers
- CPU scheduling information
 - Priorities, scheduling queue pointers
- Memory-management information
 - Memory allocated to the process
- Accounting information
 - CPU used, clock time elapsed since start, time limits
- I/O status information
 - I/O devices allocated to process, list of open files



15



CSE

Context Switch

16



Process Scheduling

- **CPU Scheduler** (short-term scheduler)
 - Selects which processes should be executed next and allocates the CPU
 - Invoked very frequently (milliseconds)
- Processes can be described as either
 - **I/O-bound**
 - Spends more time doing I/O than computations; many short CPU bursts
 - **CPU-bound**
 - Spends more time doing computations; few very long CPU bursts

17



Process Scheduling

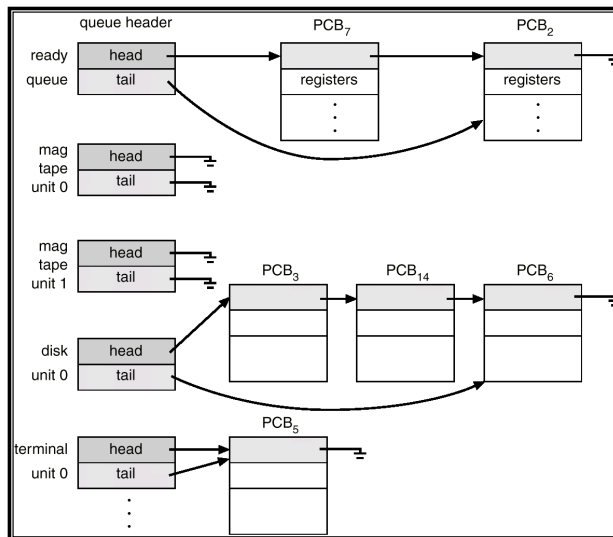
- **Job Scheduler** (long-term scheduler)
 - Selects which processes to be brought into ready queue
 - Invoked very infrequently (seconds, minutes)
 - Controls **degree of multiprogramming**, the max number of processes accommodate efficiently
- Maintains scheduling queues of processes
 - **Job queue**
 - Set of all PCBs in the system
 - **Ready queue**
 - Set of all processes residing in main memory, ready and waiting to execute
 - **Device queues**
 - Set of processes waiting for an I/O device
 - Processes migrate among the various queues



18



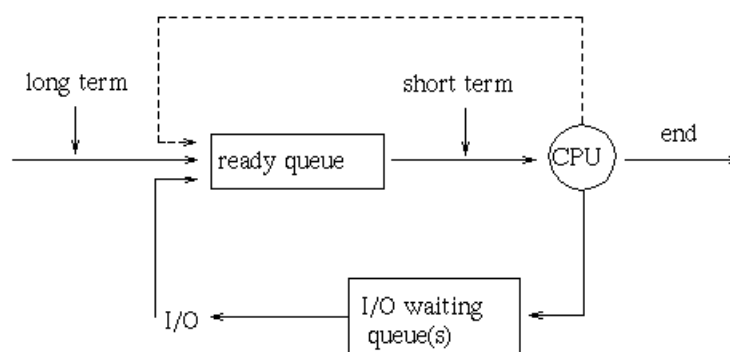
Ready Queue and I/O Device Queues



19



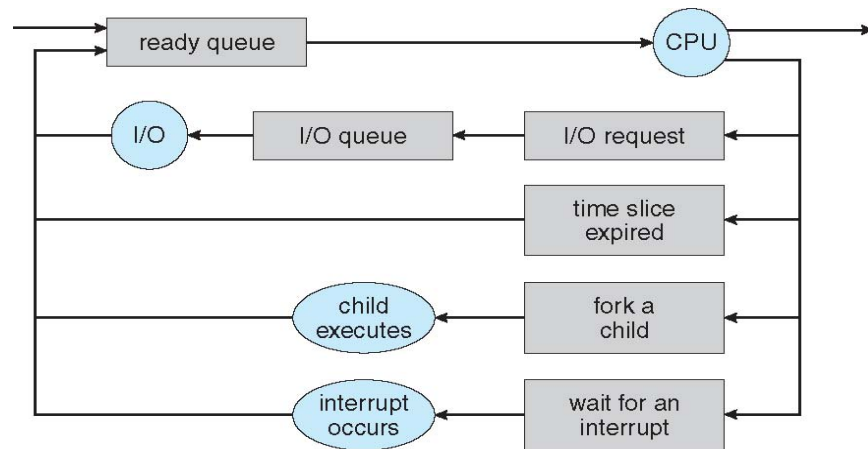
Long- and Short-Term Schedulers



20



Process Scheduling Representation



Queueing diagram represents queues, resource flows

21



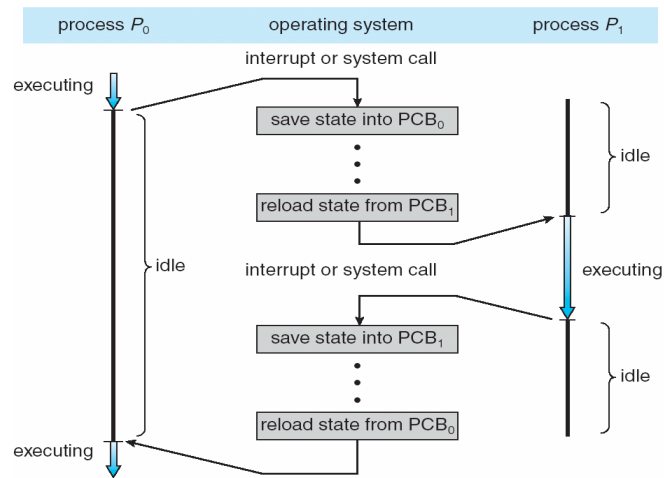
Context Switch

- When CPU switches to another process, the OS must
 - Save the PCB of old process being swapped out
 - Select new process to be swapped in
 - Load the PCB of the new process being swapped in
- Context-switch time is **overhead**
 - The system does no useful work while switching
 - Typical time about 1 μ sec
 - The more complex the OS and the PCB, the longer the context switch
- Time dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU, resulting in multiple contexts loaded at once

22



Context Switch



23



CSE

Process Execution

24



The `exec()` System Call

- Calling one of the `exec()` family of system calls
 - Terminates the currently running program, and
 - Replaces it with a new specified program that starts executing in its `main()` function
- The process ID does not change across an `exec` because a new process is *not* created
- `exec()` merely replaces the current process (its text, data, heap, and stack segments) with a brand new program from disk

25



The `exec()` Family

- There are 6 versions of the `exec` function, and they all do about the same thing:
 - Main difference is how parameters are passed

```
#include <unistd.h>
int execlp(char *file, char *arg0,
           char *arg1, ..., (char *)0);

execlp("sort", "sort", "-n",
       "city", (char *)0);
```

Same as "sort -n city"

26



The exec () Family

```
int execl(const char *path, const char *arg, ... );
    - execl takes full path name of command and variable length of
      arguments terminated by NULL
    execl("/bin/ls", "/bin/ls", "-r", "-t", "-l", NULL);

int execlp(const char *file, const char *arg, ... );
    - execlp will try to find the command from $PATH, so full path to
      command not needed
    execlp("ls", "ls", "-r", "-t", "-l", NULL);

int execlenv(const char *path, const char *arg, ...,
             char *const envp[] );
    - execlenv uses an argument list and environment variables
    char *env[] = {"PATH=/bin", NULL};
    execlenv("child.exe", "child", "arg1", "arg2", NULL,
    env);
```

27



The exec () Family

```
int execlv(const char *path, char *const argv[]);
    - execlv is the equivalent of execl, except that the arguments are
      passed in as a NULL terminated array
    char *args[] = {"/bin/ls", "-r", "-t", "-l", NULL};
    execl("/bin/ls", args);

int execlvp(const char *file, char *const argv[]);
    - execlvp is the equivalent of execlp, except that the arguments are
      passed in as a NULL terminated array
    char *args[] = {"ls", "-r", "-t", "-l", NULL};
    execlvp("ls", args);

int execlve(const char *filename, char *const
            argv [], char *const envp[] );
```

28



The `exec()` Family

- All six return `-1` on error, but no return on success
- Accept either a pathname or filename argument
- Command-line arguments are specified as separate arguments or we have to build an array of pointers to the arguments and pass the address of the array

29




`execlp` Example

```
#include <stdio.h>
#include <unistd.h>
```

```
int main()
{
```

```
    char *cmd[] = {"who", "ls", "date"};
    int i;
    printf("0=who 1=ls 2=date : ");
    scanf("%d", &i);
    execlp(cmd[i], cmd[i], (char *)0 );
    printf("execlp failed\n");
    return 0;
}
```



`printf()` not executed
unless there is a problem
with `execlp()`

30



CSE

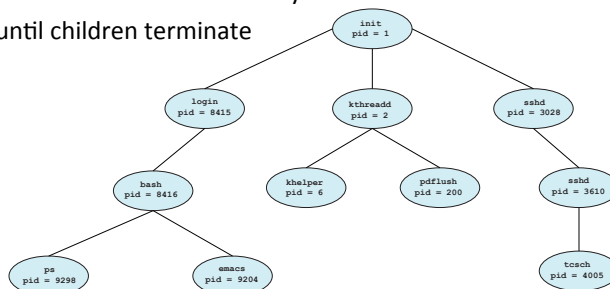
Process Creation

31



Process Creation

- A **parent** process creates **child** processes, which create other processes, forming a tree of processes
 - Generally, processes identified and managed via a process identifier (PID)
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate



32



The `fork()` System Call

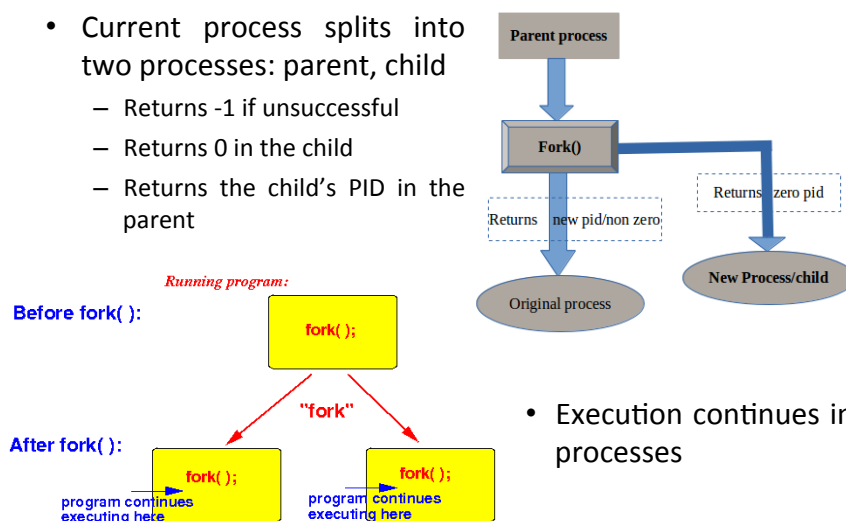
- Creation of a *new* process accomplished using the `fork()` system call
 - `fork()` creates a **child process** by making an *exact copy* of the parent process *and* then starts the process concurrently
 - `fork()` system call is unique
 - Called once, but returns twice with the parent receiving the child's unique PID and the child receiving 0 as the return value from `fork()`
 - The process that initiates the `fork()` becomes the parent process of the newly created child process
 - The child process inherits a copy of the parent's memory space, but they do *not* share the same memory as both parent and child processes will execute in their own environments

33



The `fork()` System Call

- Current process splits into two processes: parent, child
 - Returns -1 if unsuccessful
 - Returns 0 in the child
 - Returns the child's PID in the parent



34



The `fork()` System Call


- There are two uses for `fork()`:
 - When a process wants to duplicate itself so that the parent and child can each execute different sections of code at the same time
 - This is common for network servers – the parent waits for a service request from a client
 - When the request arrives, the parent calls `fork()` and lets the child handle the request
 - Parent goes back to waiting for the next service request to arrive
 - When a process wants to execute a different program
 - This is common for shells
 - In this case, the child typically does an `exec()` right after it returns from the `fork`

35



The `fork()` System Call

- The two main reasons for `fork` to fail are
 - If there are already too many processes in the system (which usually means something else is wrong)
 - If the total number of processes for this real user ID exceeds the system's limit



36



fork() Example

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main()
{
    pid_t pid;          /* could be int */
    int i;
    pid = fork();
    printf("PID=%d\n", pid);
    if( pid > 0 )
    {
        /* parent */
        for( i=0; i < 10; i++ )
            printf("\t\t\tPARENT %d\n", i);
    }
    else
    {
        /* child */
        for( i=0; i < 10; i++ )
            printf("CHILD %d\n", i );
    }
    return 0;
}
```

37



fork() Example

- Notes for code on previous slide
 - `i` is copied between parent and child
 - The switching between the parent and child depends on many factors:
 - Machine load, system process scheduling
 - I/O buffering affects amount of output shown
 - Output interleaving is *nondeterministic*
 - Cannot determine exact output by looking at code

38

**CSE**

Process Synchronization

39



`wait()` and `waitpid()` System Calls

- The `wait()` and `waitpid()` system calls force the parent process to suspend execution until the child process has completed
 - `waitpid()` waits for a specific child process identified by its PID while `wait()` simply waits for the first child process to terminate (if the parent has more than one child process)
 - Both return the PID of the terminated process if successful
 - Or `-1` if an error occurred (usually means no child exists to wait on)
 - Once the child process has terminated, the parent process resumes execution

40



wait() Function

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t wait(int *statloc);
```

- **statloc** can be (int *)0 or a variable which will be bound to status information about the child
- A process that calls **wait()** can
 - **suspend** (block) if all of its children are still running
 - return immediately with the termination status of a child
 - return immediately with an error if there are no child processes.


41



wait() Example

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

int main() {
    char *cmd[] = {"who", "ls", "date"};
    int i;
    while( 1 ) {
        printf("0=who 1=ls 2=date : ");
        scanf("%d", &i);
        if(fork() == 0) {
            /* child */
            execlp( cmd[i], cmd[i], (char *)0 );
            printf("execlp failed\n");
            exit(1);
        }
        else {
            /* parent */
            wait( (int *)0 );
            printf("child finished\n");
        }
    } /* while */
} /* main */
```




42



The fork-exec Model

- Common use of the `fork()` and `exec()` system calls centers around being able to run another program in parallel without *terminating* the current process
 - Parent process uses `fork()` to create the child process, which will then use the `exec()` family of system calls to run the desired program while the parent waits on the child to terminate



43



CSE

Process Termination

44



Process Termination

- A process may be terminated
 - When it executes its last statement and asks the operating system to delete it
 - By calling the `exit()` system call in the `<stdlib.h>` library
- A parent may terminate execution of child processes using the `abort()` system call if
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - Parent is exiting
 - Some operating systems do not allow child to continue if its parent terminates
- When a child process terminates, a **SIGCHLD** signal is sent to the parent

45



Zombies and Orphans

- A child process whose parent has terminated is referred to as an **orphan**
 - Child is still executing, but parent has terminated
 - Some process is needed to query the child's exit status
- When a child exits and its parent is not currently waiting (i.e., executing a `wait()`), it becomes a **zombie**
 - A zombie is not really a process (since it terminated), but the system still has an entry in the process table for the non-existing child process
- When a parent terminates, any orphans and zombies are adopted by the **init** process (PID 1) of the system

46



Process Identification

47



Process Identification

- `pid = getpid();`
 - Returns its own process id
- `pid = getppid();`
 - Returns parent process id
- `uid = getuid();`
 - Returns real user ID of own process
- `newpg = setpgrp();`
 - Sets process group of own process to itself
- `pgid = getpgrp();`
 - Returns the process group ID of own process

48



Process and Group ID's Example

```
#include <stdio.h>
#include <sys/types.h>
#include <stdlib.h>
#include <unistd.h>
int main() {
    pid_t cpid, pid, pgid, cpgid; //process id's and process groups
    cpid = fork();
    if (cpid == 0) {
        /* CHILD */
        //set process group to itself
        setpgid(); //<-----needed to disambiguate runtime process!
        //print the pid, and pgid of child from child
        pid = getpid();
        pgid = getpgid();
        printf("Child:          pid:%d pgid:%d\n", pid, pgid);
    }
    else if (cpid > 0) {
        /* PARENT */
        //set the process group of child
        setpgid(cpid, cpid); //<-----needed to disambiguate runtime process
        //print the pid, and pgid of parent
        pid = getpid();
        pgid = getpgid();
        printf("Parent:          pid:%d pgid: %d \n", pid, pgid);
        //print the pid, and pgid of child from parent
        cpgid = getpgid(cpid);
        printf("Parent: Child's pid:%d pgid:%d\n", cpid, cpgid);
    }
    else {
        /*ERROR*/
        perror("fork");
        exit(1);
    }
    return 0;
}
```

With this, it will not matter which runs first, parent or child, as the result will be the same – the child placed in the appropriate process group

49



CSE

Process Data

50



Process Data

- Since a child process is a copy of the parent, it has *copies* of the parent's data
- A change to a variable in the child will *not* change that variable in the parent

51



Process Data Example

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int glbvar = 6;
int main() {
    int locvar = 88;
    pid_t pid;
    printf("Before fork()\n");
    if( (pid = fork()) == 0 ) {
        /* child */
        glbvar++;
        locvar++;
    }
    else if ( pid > 0 ) {
        /* parent */
        sleep(2);
    }
    else
        perror("fork error");
    printf("pid=%d, glbvar=%d, locvar=%d\n", getpid(), glbvar, locvar);
    return 0;
} /* end main */
```

52



Inherited Data and File Descriptors

- A forked child has instances of current values of the variables and open file descriptors
- Variables
 - Passed by value (i.e., a copy)
- Read/write pointers for a file
 - Passed by reference

53



Process File Descriptors

- A child and parent have copies of the file descriptors, but the R-W pointer is maintained by the system
 - The R-W pointer is **shared**
- This means that a `read()` or `write()` in one process will affect the other process since the R-W pointer is changed

54



Process File Descriptors Example

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <fcntl.h>
void printpos(char *msg, int fd);
int main() {
    int fd; /* file descriptor */
    pid_t pid;
    char buf[10]; /* for file data */
    if ((fd=open("file1", O_RDONLY)) < 0)
        perror("open");
    read(fd, buf, 10); /* move R-W ptr */
    printpos("Before fork", fd);
    if( (pid = fork()) == 0 ) {
        /* child */
        printpos("Child before read", fd);
        read(fd, buf, 10);
        printpos("Child after read", fd);
    }
    else if( pid > 0 ) {
        /* parent */
        wait((int *)0);
        printpos("Parent after wait", fd);
    }
    else
        perror("fork");
}
```

55



Process File Descriptors Example

```
/* Print position in file */
void printpos(char *msg, int fd) {
    long int pos;
    if( (pos = lseek( fd, 0L, SEEK_CUR) ) < 0L )
        perror("lseek");
    printf("%s: %ld\n", msg, pos);
}
```

```
$ ./a.out
Before fork: 10
Child before read: 10
Child after read: 20
Parent after wait: 20
```

what's happened?

56