

# CSCE 3600

## Principles of Systems Programming

### Interprocess Communication Part 2

University of North Texas




## What are Sockets?

- Sockets are an **extension of pipes**, with the advantages that processes don't need to be related, or even on the same machine
  - Method for accomplishing interprocess communication
  - Bidirectional communication
- A socket is like the end point of a pipe – in fact, the Linux kernel implements pipes as a pair of sockets
- Two (or more) sockets must be connected before they can be used to transfer data



## Socket Client-Server Model

- IPC using sockets described as **client-server model**
  - One process is usually called the server
  - A server process is usually responsible for satisfying requests made by other processes called clients
  - A server usually has a well-known address (e.g., IP address or pathname)



3



## Socket Attributes

- Sockets are characterized by three attributes
    - Domain
    - Type
    - Protocol
- Socket Protocol
    - Determined by socket type and domain
    - Default protocol is 0
- For communication between processes, sockets can be implemented in the following domains
    - UNIX (e.g., `AF_UNIX`)
      - Processes are on the same machine
    - INET (e.g., `AF_INET`)
      - Each process is on a different machine (i.e., requires a network interface device)

4



## Socket Attributes

---

- Three main types of sockets:
  - **SOCK\_STREAM** (TCP sockets)
    - Provides a connection-oriented, sequenced, reliable, and bidirectional network communication service
    - E.g., telnet, ssh, http
  - **SOCK\_DGRAM** (UDP sockets)
    - Provides a connectionless, unreliable, best-effort network communication service
    - E.g., streaming audio/video, IP telephony
  - **SOCK\_RAW**
    - Allows direct access to other layer protocols such as IP, ICMP, or IGMP

5

**CSE**

## Unix Domain Stream Sockets

6



## UNIX Domain Sockets

- Socket for communicating with another process on the **same machine only**
  - Provides an optimization since no network overhead
- Uses `sockaddr_un` structure instead of `sockaddr_in`

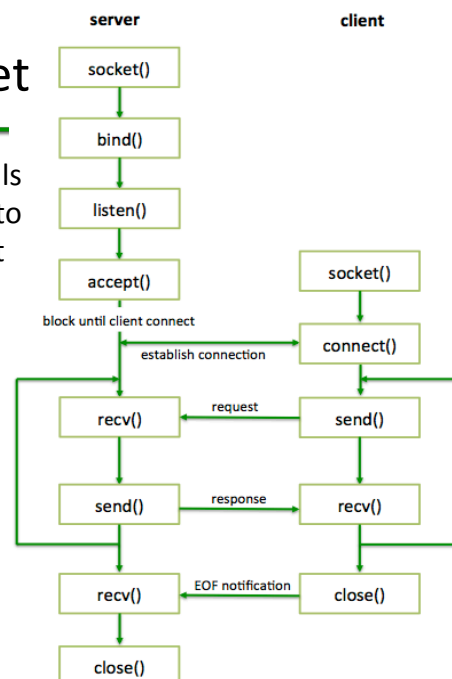
```
#include <sys/un.h>
struct sockaddr_un
{
    sa_family_t sun_family;    /* AF_UNIX */
    char sun_path[108];        /* pathname */
};
```

7



## Stream Socket

- Sequence of function calls for a client and server to implement a stream socket



8



## Steps in Server Process

---

1. Call `socket()` with proper arguments to create the socket
2. Call `bind()` to bind the socket to an address (in our case, it is just a pathname) in the UNIX domain
3. Call `listen()` to instruct the socket to listen for incoming connections from client programs
4. Call `accept()` to accept a connection from a client
5. Handle the connection and loop back to `accept()`
6. Close the connection

9



## Steps in Client Process

---

1. Call `socket()` to get a UNIX domain socket to communicate through
2. Set up a `struct sockaddr_un` with the remote address (where the server is listening) and call `connect()` with that as an argument
3. Assuming no errors, you are connected to the remote side
  - Use `send()` and `recv()` to communicate

10



## Creating the Socket

```
#include <sys/types.h>
#include <sys/socket.h>
int socket(int domain, int type, int protocol);
```

- Creates an endpoint for communication (i.e., a socket) and returns a file descriptor
- Common domains (or **address family**)
  - `AF_UNIX` Unix domain sockets
  - `AF_INET` IPv4 Internet sockets
  - `AF_INET6` IPv6 Internet sockets
- Type
  - `SOCK_STREAM`, `SOCK_DGRAM`, `SOCK_RAW`
- Protocol
  - Set to 0 (chosen by OS) except for RAW sockets

11



## Bind to a Name (Server Side)

```
#include <sys/types.h>
#include <sys/socket.h>
int bind(int sockfd, const struct sockaddr *serveraddr,
        socklen_t addrlen);
```

- Binds a name to the socket (i.e., reserves a port)
- `sockfd` is the socket file descriptor returned by `socket`
- `serveraddr` is the IP address and port of the machine (address usually set to `INADDR_ANY` – chooses a local address)
- `addrlen` is the size (in bytes) of the structure
- Returns 0 on success or `-1` if an error occurs
- When socket is bound, new special socket-type file (type "s") corresponding to `sun_path` is created
- This file is *not* automatically deleted, so need to `unlink` it
  - If `bind` finds the file already exists, it will fail

12



## Set Up Queue (Server Side)

```
#include <sys/types.h>
#include <sys/socket.h>
int listen(int sockfd, int backlog);
```

- Listens for connections on a socket
  - After calling `listen()`, a socket is ready to accept connections
- Prepares a queue in the kernel where partially completed connections wait to be accepted
- Many client requests may arrive
  - Server cannot handle them all at the same time
  - Server could reject the requests or let them wait
  - `backlog` is the maximum number of partially completed connections that the kernel should queue
- What if too many clients arrive?
  - Some requests don't get through ... makes no promises
  - And the client(s) can always try again!

13



## Establish Connection (Server Side)

```
#include <sys/types.h>
#include <sys/socket.h>
int accept(int sockfd, struct sockaddr* cliaddr,
           socklen_t * addrlen);
```

- Accepts a connection on the socket and returns a new file descriptor that refers to the connection with the client that will be used for reads and writes on the connection
  - Blocks waiting for a connection (from the queue)
- `sockfd` is the listening socket
- `cliaddr` is the address of the client

14



## Establish Connection (Client Side)

```
#include <sys/types.h>
#include <sys/socket.h>
int connect(int sockfd, const struct sockaddr*
            servaddr, socklen_t * addrlen);
```

- Initiates a connection on the socket, where the kernel will select a source IP address and dynamic port
- Returns 0 on success or -1 on failure

15



## Sending Data

```
#include <sys/types.h>
#include <sys/socket.h>
ssize_t send(int sockfd, const void *buf, size_t len,
             int flags);
```

- Alternative to `write` that sends a message on the socket
- `sockfd` is socket descriptor for open and connected socket
- `buf` is a buffer of information to send
- `len` is the size of the buffer in bytes
- `flags`
  - Bitwise OR of zero or more flags (e.g., `MSG_EOR`, `MSG_OOB`, etc.)
- Calls are **blocking** (i.e., returns only after data is sent)

16





## Receiving Data

```
#include <sys/types.h>
#include <sys/socket.h>
ssize_t recv(int sockfd, void *buf, size_t len, int
             flags);
```

- Alternative to `read` that receives a message on the socket
- `sockfd` is socket descriptor for open and connected socket
- `buf` is initially empty to store the data
- `len` is the size of the buffer in bytes
- `flags`
  - Bitwise OR of zero or more flags (e.g., `MSG_EOR`, `MSG_OOB`, etc.)
- Calls are **blocking** (returns only after data is received)

17



## Close and Shutdown

```
#include <unistd.h>
int close(sockfd);
```

- Closes the file descriptor when done
- Will not deallocate the socket until we close the **last** descriptor that references it (we may have several)

```
#include <sys/socket.h>
int shutdown(int sockfd, int how);
```

- Forces a full or partial closure of a socket
- `sockfd` is a socket descriptor
- `how` can be `SHUT_RD`, `SHUT_WR`, or `SHUT_RDWR`

18

**CSE**

## Socket Pair

19



## Socket Pipe

- What if you wanted a `pipe()`, but you wanted to use a single pipe to send and receive data from both sides?
  - Since most pipes are unidirectional (with exceptions in SYSTEM V), you cannot do it!
- UNIX domain sockets can handle bi-directional data
  - But, you would have to set up all that code with `listen()` and `connect()` with everything just to pass data both ways
- However, you can use `socketpair()` that returns a pair of already connected sockets
  - No extra work needed and you can immediately use the socket descriptors for interprocess communication (only available in UNIX domain socket)

20



## socketpair Function

---

```
#include <sys/socket.h>
int socketpair(int family, int type, int protocol, int
    sockfd[2]);
```

- Creates two **unnamed** sockets that are already connected
- They are **full-duplex** (i.e., data can go in each direction)
- Also called **stream pipes** if type is **SOCK\_STREAM**
- Returns 0 on success, -1 on error
- Family must be **AF\_LOCAL** or **AF\_UNIX**
- Type is **SOCK\_DGRAM** or **SOCK\_STREAM**
- Protocol must be 0

21

**CSE**


---

## Internet Stream Sockets

22



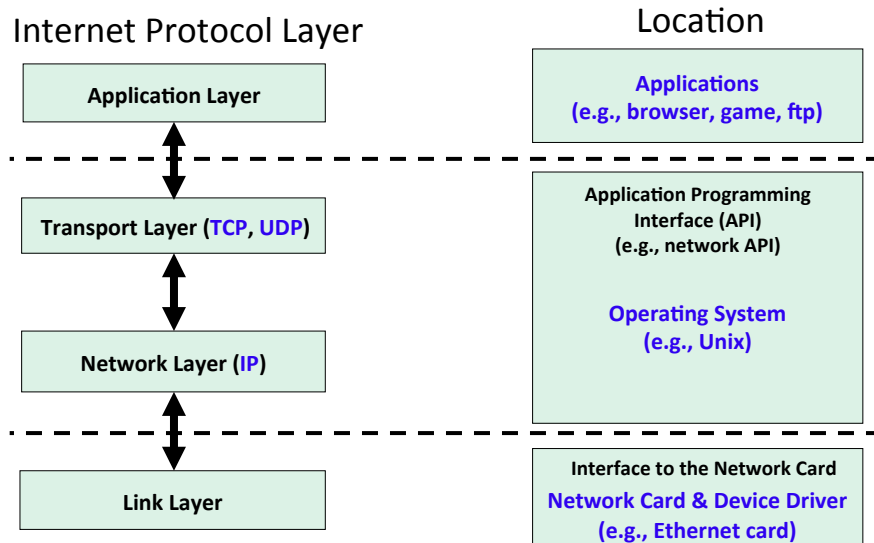
## Layers of the IP Protocol Suite



23



## Protocol Suite Location




24



## Internet Sockets

- Sockets provide a mechanism to communicate between computers across a **network**
- Internet Sockets
  - AF\_INET – the addresses are IPv4 addresses
  - AF\_INET6 – the addresses are IPv6 addresses



25



## Byte Ordering of Integers

- Different CPU architectures have different byte ordering
- **Little-Endian**
  - Stores the least significant byte in the smallest address
- **Big-Endian**
  - Stores the most significant byte in the smallest address
- In order to connect to a remote computer and use a socket, we need to use its address
  - In fact, Linux is little-endian, but TCP/IP uses big-endian byte ordering

26



## Byte Ordering

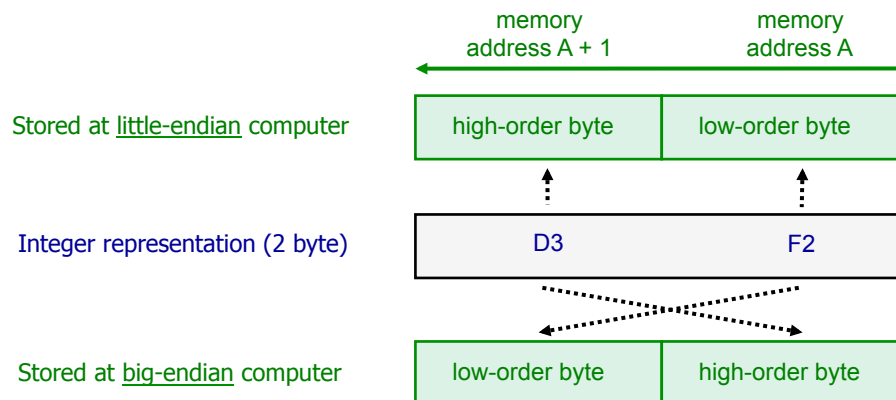
- Suppose we wanted to store the integer 305,441,741
  - Converts to  $1234ABCD_{16}$

Little-Endian Byte Order		Big-Endian Byte Order	
Address	Value	Address	Value
Low Memory Address 0x8000	0xCD	Low Memory Address 0x8000	0x12
0x8001	0xAB	0x8001	0x34
0x8002	0x34	0x8002	0xAB
High Memory Address 0x8003	0x12	High Memory Address 0x8003	0xCD

27



## Byte Ordering of Integers

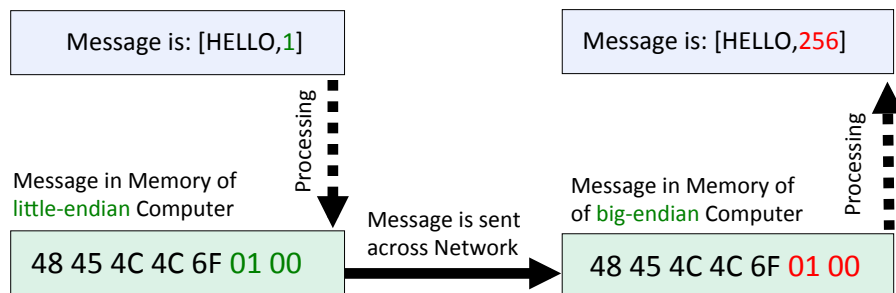


28



## Byte Ordering Problem

- What would happen if two computers with different integer byte ordering communicate?
  - Nothing... if they do not exchange integers!
  - But... if they exchange integers, they would get the wrong order of bytes, and therefore, the wrong value!



29



## Byte Ordering Solution

- There are two solutions if computers with different byte ordering system want to communicate
  - They must know the kind of architecture of the sending computer
    - Bad solution, it has not been implemented
  - Introduction of a network byte order. The functions are:
 

```
uint16_t htons(uint16_t host16bitvalue)
uint32_t htonl(uint32_t host32bitvalue)
uint16_t ntohs(uint16_t net16bitvalue)
uint32_t ntohl(uint32_t net32bitvalue)
```
- Use for all integers (short and long) that are sent across the network
  - Including port numbers and IP addresses

30



## Naming and Addressing


- Host name
  - Identifies a single host
  - Variable length string (e.g., www.unt.edu, cse01)
  - Is mapped to one or more IP addresses
- IP Address (IPv4)
  - Written as dotted octets (e.g., 10.0.0.1)
  - 32 bits – not a number, but often needs to be converted to a 32-bit number to use
- Port number
  - Identifies a process on a host
  - 16 bit number

31



## Addresses, Ports, and Sockets

- Like apartments and mailboxes
  - You are the application
  - Street address of your apartment building is the **IP address**
  - Your mailbox is the **port**
  - The post-office is the network
  - The **socket** is the key that gives you access to the right mailbox
- How do you choose which port a socket connects to?




32





## Port Numbers




- You can find a list of all IANA registered ports at:
  - <http://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml>

33



## Addresses, Ports, and Sockets

- Choose a port number that is registered for general use, from 1024 to 49151
  - Do not use ports 0 to 1023. These ports are **reserved** (must be **root**) for use by the Internet Assigned Numbers Authority (IANA)
  - Avoid using ports 49152 through 65535. These are dynamic ports that operating systems use randomly
    - If you choose one of these ports, you risk a potential port conflict



34



## IP Address Data Structure

- The `sockaddr_in` structure has four parts

```
struct sockaddr_in
{
    short int          sin_family; // address family
    unsigned short int sin_port;   // port number
    struct in_addr     sin_addr;   // Internet address
    unsigned char      sin_zero[8]; // set to 0, padding
};

struct in_addr
{
    unsigned long      s_addr;     // 4 bytes
};
```

35



## IP Address Data Structure

- Declare address structure
 

```
struct sockaddr_in sockAdd;
```
- Set family
 

```
sockAdd.sin_family = AF_INET;
```
- Set IP address (2 ways)
 

```
// specify address to listen to
inet_pton(AF_INET, "127.0.0.1", &sockAdd.sin_addr.s_addr)
// listen to any local address
sockAdd.sin_addr.s_addr = htonl(INADDR_ANY)
```
- Set port
 

```
sockAdd.sin_port = htons(9999);
```

36



## Network Addressing Library Routines

```
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
in_addr_t inet_addr(const char* string)
```

- Converts an IP address in IPv4 numbers-and-dots notation into binary form in network byte order
- If string not contain legitimate Internet address, returns value `INADDR_NONE`

```
#include <unistd.h>
int gethostname(char* name, int nameLen)
```

- Gets the null-terminated hostname as a character array along with its length in bytes

37



## Network Addressing Library Routines

```
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
int inet_aton(const char *cp, struct in_addr *inp);
```

- Converts an IP address in IPv4 numbers-and-dots notation into binary form in network byte order and stores it in an `in_addr` struct

```
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
char *inet_ntoa(struct in_addr in);
```

- Converts an IP address in the `in_addr` struct in network byte order into a string in IPv4 dotted-decimal notation

38



## Other Library Routines

---

```
#include <arpa/inet.h>
const char *inet_ntop(int af, const void *src, char
    *dst, socklen_t cnt);
```

- Converts a network address, either IPv4 or IPv6, from its binary to text form
- 

```
#include <arpa/inet.h>
int inet_pton(int af, const char *src, void *dst);
```

- Converts a network address, either IPv4 or IPv6, from its text to binary form

39

**CSE**

---

## Datagram Sockets

40



## TCP vs. UDP

---

- TCP
  - Reliable byte stream service
  - Different ways to build clients and servers
    - Some problems when building clients/servers with TCP
      - How to get around blocking I/O
      - Data format conversion
    - Basic assumption: whatever sent will eventually be received!!
- UDP
  - Unreliable datagram service
    - Data may get lost – application may need to deal with more details in the communication

41



## Why UDP?

---

- Applications that do not need 100% reliability communication
  - E.g., VoIP, video stream, DNS servers
- Applications care a lot about performance
  - High performance computing (TCP cares too much about other things than performance)
- Applications that need multicast or broadcast
  - TCP only supports point to point communication)

42



## UDP Client/Server

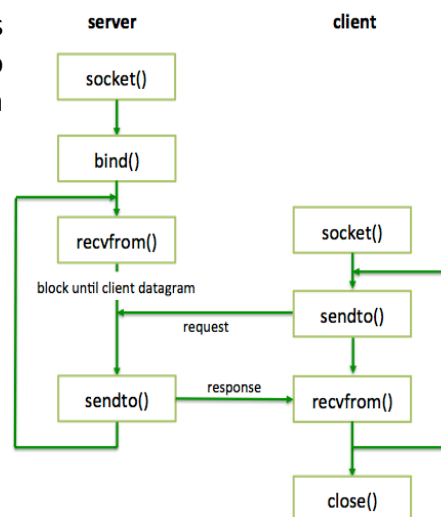
- Typical UDP client
  - Client does *not* establish a connection with the server
  - Client sends a datagram to the server using `sendto` function
- Typical UDP server
  - Does *not* accept a connection from a client
  - Server calls `recvfrom` function that waits until data arrives from some client

43



## Datagram Socket

- Sequence of function calls for a client and server to implement a datagram socket
  - No “handshake”
  - No simultaneous close



44



## Receiving a Message

```
#include<sys/socket.h>
ssize_t recvfrom(int sockfd, void *buff, size_t nbytes,
    int flags, struct sockaddr *from, socklen_t
    *addrlen);
```

- Receives a message on the socket
- If `recvfrom` is successful, the number of bytes read is returned, -1 on error

45



## Sending a Message

```
#include <sys/socket.h>
ssize_t sendto(int sockfd, const void *buff, size_t
    nbytes, int flags, const struct sockaddr *to,
    socklen_t addrlen);
```

- Sends a message on the socket
- If `sendto` is successful, the number of bytes written, even if 0 bytes are written, is returned, -1 on error

46



CSE

## I/O Multiplexing

47



## Dealing with Blocking Calls

- Many functions block
  - `accept` until a connection comes in
  - `connect` until the connection is established
  - `recv`, `recvfrom` until a packet (of data) is received
  - `send`, `sendto` until data is pushed into socket's buffer
- For simple programs, this is fine
- What about complex connection routines?
  - Multiple connections
  - Simultaneous sends and receives
  - Simultaneously doing non-networking processing

48





## Dealing with Blocking Calls

- Options
  - Create multi-process or multi-threaded code
  - Turn off blocking feature (e.g., use `fcntl()` file-descriptor control system call)
  - Use the `select()` function
- What does `select()` do?
  - Can be permanent blocking, time-limited blocking, or non-blocking
  - Input is a set of file descriptors
  - Output is information on the file-descriptors' status
  - Therefore, can identify sockets that are "ready for use" so that calls involving that socket will return immediately

49



## I/O Blocking

```
socket();
bind();
listen();
while
    accept();
    recv();
    send();
```

- Simple server has **blocking** problem
  - Suppose 5 connections accepted
  - Suppose next `accept()` blocks
  - Other connections cannot send and receive
  - Cannot get keyboard input either

50



## I/O Multiplexing

- `select()` waits on multiple file descriptors and timeout
- Returns when any file descriptor
  - Is ready to be read, or
  - Written, or
  - Indicates an error, or
  - Timeout exceeded
- Advantages
  - Simple
  - Application **does not consume CPU cycles while waiting**
- Disadvantages
  - **Does not scale to large number** of file descriptors

```
socket();
bind();
listen();
while
    accept();
    recv();
    send();
```

wait for `select()`

51



## select Function Call

```
#include <sys/types.h>
#include <sys/socket.h>
int status = select(nfds, &readfds, &writefds,
    &exceptfds, &timeout);
```

- Provided a list of file descriptors to check the status if they are **read-ready**, **write-ready**, or have registered an **exception**
- Returns the number of ready objects, -1 if error
- `nfds` is 1 + largest file descriptor to check
- Uses bit-vector structure called `fd_set` to manage list of file descriptors
  - `readfds` list of descriptors to check if read-ready
  - `writefds` list of descriptors to check if write-ready
  - `exceptfds` list of descriptors to check if an exception is registered
- `timeout` passed to indicate time after which `select` returns, even if no ready file descriptors

52



## To Be Used With select

---

- `select` uses a structure `struct fd_set`
  - It is just a bit-vector
  - If bit `i` is set in [`readfds`, `writfds`, `exceptfds`], `select` will check if file descriptor (i.e. socket) `i` is ready for [reading, writing, exception]
- Before calling `select`
  - `FD_ZERO(&fdvar)` clears the structure
  - `FD_SET(i, &fdvar)` to check file descriptor `i`
- After calling `select`
  - `int FD_ISSET(i, &fdvar)` boolean returns TRUE iff `i` is “ready”